

GenMol UDLM engineering exploration

engineering-v5-temperature | complete | 2026-09-07T09:50:06.376168+00:00

Engineering pilots only. Superiority has not been established.

Engineering pilots only: no superiority claim, final candidate lock, or final-seed evaluation.

The local MDLM model completed 50,000 updates; each UDLM arm starts from its EMA weights and adds 1,000 updates. Training compute is not matched.

The scheduler and FiLM conditioner were selected on E denoising loss and then shared by R/S/E. This is not independent optimization of each method.

Small samples and configuration selection create sampling uncertainty and selection bias; a higher pilot mean is not evidence of superiority.

Strict decoding uses fix=False. Released-compatible decoding uses SAFE repair and largest-component selection; report both because repairs can hide malformed outputs.

Validity and quality divide by requested samples. Uniqueness divides by valid samples. Quality counts unique molecules with QED $\geq$ 0.6 and SA $\leq$ 4. Diversity is the released PyTDC fingerprint metric over unique valid molecules.

Means and sample standard deviations average seeds equally within one configuration. Diversity is not pooled across seeds. Failed, pending, and invalid runs are excluded from means and remain disclosed.

Paper seeds are undisclosed and hardware differs. Local timings are descriptive; no speed claim relative to the paper is made.

Reference means (context only)

Reference	validity	uniqueness	quality	diversity
local_mdmlm_50000	1.000000	0.998667	0.858000	0.823021
paper_genmol_v1_table1	1.000000	0.997000	0.846000	0.818000

Per-configuration mean, equal seed weights

released_comparable

Config	Seeds complete	validity	uniqueness	quality	diversity
e_t050_p100	2/2	0.9766 +/- 0.0110	1.0000 +/- 0.0000	0.4531 +/- 0.0663	0.8921 +/- 0.0055
e_t070_p100	2/2	0.9844 +/- 0.0000	0.9921 +/- 0.0112	0.4531 +/- 0.1105	0.8994 +/- 0.0028
e_t085_p100	2/2	0.9453 +/- 0.0110	0.9668 +/- 0.0238	0.3984 +/- 0.0331	0.9038 +/- 0.0054
e_t100_p100	2/2	0.9375 +/- 0.0442	0.9914 +/- 0.0122	0.3594 +/- 0.0663	0.9201 +/- 0.0115

Config	Seeds complete	validity	uniqueness	quality	diversity
s_t050_p100	2/2	0.9766 +/- 0.0110	0.9919 +/- 0.0114	0.4141 +/- 0.1436	0.8967 +/- 0.0043
s_t070_p100	2/2	0.9766 +/- 0.0110	1.0000 +/- 0.0000	0.4219 +/- 0.1989	0.8983 +/- 0.0064
s_t085_p100	2/2	0.9609 +/- 0.0110	1.0000 +/- 0.0000	0.5078 +/- 0.0552	0.9045 +/- 0.0044
s_t100_p100	2/2	0.9297 +/- 0.0773	0.9841 +/- 0.0224	0.3203 +/- 0.0110	0.9242 +/- 0.0053
r_t050_p100	2/2	0.9922 +/- 0.0110	1.0000 +/- 0.0000	0.4844 +/- 0.0442	0.8809 +/- 0.0010
r_t070_p100	2/2	0.9766 +/- 0.0331	1.0000 +/- 0.0000	0.4453 +/- 0.0110	0.8942 +/- 0.0019
r_t085_p100	2/2	0.9688 +/- 0.0221	0.9921 +/- 0.0112	0.3594 +/- 0.0663	0.9114 +/- 0.0020
r_t100_p100	2/2	0.9375 +/- 0.0000	0.9917 +/- 0.0118	0.3047 +/- 0.0110	0.9110 +/- 0.0051

strict

Config	Seeds complete	validity	uniqueness	quality	diversity
e_t050_p100	2/2	0.4766 +/- 0.0110	1.0000 +/- 0.0000	0.2266 +/- 0.0552	0.8785 +/- 0.0021
e_t070_p100	2/2	0.4062 +/- 0.1105	1.0000 +/- 0.0000	0.2109 +/- 0.0994	0.8718 +/- 0.0153
e_t085_p100	2/2	0.3203 +/- 0.0331	1.0000 +/- 0.0000	0.1328 +/- 0.0331	0.8833 +/- 0.0101
e_t100_p100	2/2	0.1797 +/- 0.0331	1.0000 +/- 0.0000	0.0938 +/- 0.0442	0.8839 +/- 0.0066
s_t050_p100	2/2	0.6016 +/- 0.0110	1.0000 +/- 0.0000	0.2344 +/- 0.0884	0.8781 +/- 0.0036
s_t070_p100	2/2	0.4297 +/- 0.0110	1.0000 +/- 0.0000	0.2109 +/- 0.0110	0.8641 +/- 0.0042
s_t085_p100	2/2	0.3984 +/- 0.0552	1.0000 +/- 0.0000	0.2422 +/- 0.0331	0.8713 +/- 0.0036
s_t100_p100	2/2	0.2500 +/- 0.0000	1.0000 +/- 0.0000	0.1172 +/- 0.0110	0.8859 +/- 0.0015
r_t050_p100	2/2	0.6328 +/- 0.0552	1.0000 +/- 0.0000	0.2969 +/- 0.0442	0.8686 +/- 0.0007
r_t070_p100	2/2	0.4609 +/- 0.0110	1.0000 +/- 0.0000	0.1953 +/- 0.0552	0.8749 +/- 0.0016
r_t085_p100	2/2	0.3984 +/- 0.0331	1.0000 +/- 0.0000	0.2109 +/- 0.0552	0.8772 +/- 0.0103
r_t100_p100	2/2	0.3750 +/- 0.0663	1.0000 +/- 0.0000	0.1406 +/- 0.0221	0.8811 +/- 0.0030

Every scheduled seed

Config	Seed	Status	Requests	Raw rows	Repair Q	Strict Q	Generation s
e_t050_p100	1200	completed	64	64	0.5000	0.2656	9.71
e_t050_p100	1201	completed	64	64	0.4062	0.1875	9.17
e_t070_p100	1200	completed	64	64	0.5312	0.2812	9.55
e_t070_p100	1201	completed	64	64	0.3750	0.1406	9.09
e_t085_p100	1200	completed	64	64	0.3750	0.1094	9.58
e_t085_p100	1201	completed	64	64	0.4219	0.1562	9.24
e_t100_p100	1200	completed	64	64	0.3125	0.0625	9.46
e_t100_p100	1201	completed	64	64	0.4062	0.1250	9.70
s_t050_p100	1200	completed	64	64	0.5156	0.2969	9.42
s_t050_p100	1201	completed	64	64	0.3125	0.1719	9.15
s_t070_p100	1200	completed	64	64	0.5625	0.2188	9.41
s_t070_p100	1201	completed	64	64	0.2812	0.2031	9.05
s_t085_p100	1200	completed	64	64	0.5469	0.2656	10.77
s_t085_p100	1201	completed	64	64	0.4688	0.2188	9.01
s_t100_p100	1200	completed	64	64	0.3281	0.1094	10.08
s_t100_p100	1201	completed	64	64	0.3125	0.1250	9.10
r_t050_p100	1200	completed	64	64	0.5156	0.3281	9.99
r_t050_p100	1201	completed	64	64	0.4531	0.2656	9.02
r_t070_p100	1200	completed	64	64	0.4531	0.2344	9.30
r_t070_p100	1201	completed	64	64	0.4375	0.1562	8.96
r_t085_p100	1200	completed	64	64	0.4062	0.2500	9.33
r_t085_p100	1201	completed	64	64	0.3125	0.1719	9.02
r_t100_p100	1200	completed	64	64	0.3125	0.1250	9.35
r_t100_p100	1201	completed	64	64	0.2969	0.1562	9.06

Preserved v4 diagnostic failure

V4 remains campaign_incomplete; its failed D attempt is never reclassified or reused for ranking.

completion validation failed: Seed 1100 has benchmark artifacts that do not match the requested run or fail integrity checks: - config.effective={ 'device': 'cuda:0', 'diffusion_type': 'udlm', 'exclude_special_tokens': False, 'inference_eps': 1e-05, 'min_add_len': 40, 'model_path': '/home/aidar.alimbayev/Documents/genmolv2/run_sources/udlm_genmol_worktree/output/udlm/scaleup-w1-e-dcb271453411/checkpoints/1000.ckpt', 'num_samples': 32, 'num_steps': 128, 'prior_metadata_sha256': 'f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612', 'prior_variant': 'empirical_frequency', 'randomness': 0.0, 'raw_loo_top_p': 1.0, 'softmax_temp': 1.0}; expected { 'model_path': '/home/aidar.alimbayev/Documents/genmolv2/run_sources/udlm_genmol_worktree/output/udlm/scaleup-w1-e-dcb271453411/checkpoints/1000.ckpt', 'num_samples': 32, 'diffusion_type': 'udlm', 'softmax_temp': 1.0, 'randomness': 0.0, 'min_add_len': 40, 'num_steps': 128, 'inference_eps': 1e-05, 'exclude_special_tokens': False, 'prior_variant': 'empirical_frequency', 'prior_metadata_sha256': 'f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612', 'device': 'cuda:0' } - config.effective_sha256='278d4476e98bc0b5905e16c2158dd75a3a5a85250e4a48726640afc0e5529608'; expected '3ee30748fa7d7fcc675fde0fe1e8ff15f23fdc3374ba76eed03466ba306ec02f' Refusing to skip or relaunch into the existing directory. Inspect the artifacts or choose a fresh output root.

Configuration and artifact provenance

Protocol SHA-256: ecd4722c23debe193aeb50c034122ca20e501cf7c379e4beb6514b4d496d1a7

v5-e-t050-p100 / seed 1200: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3", "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t050-p100/seed_1200/raw_samples.csv", "sha256": "575549aadb838cc83426a13fe3bc3233922f0f4b8ca79a618a6ba630e6e9f414", "size_bytes": 20244, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}}

v5-e-t050-p100 / seed 1201: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5", "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t050-p100/seed_1201/raw_samples.csv", "sha256": "8894501102bb1dc32ae64feec0d3e932dd7a24a364ae4ca6dc74d76db8e7febb", "size_bytes": 20232, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}}

v5-e-t070-p100 / seed 1200: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3", "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t070-p100/seed_1200/raw_samples.csv", "sha256": "a0832e0b5489c32d25e9732b24e9160b32c48c485b39406f45d3a54c0ecc7cf8", "size_bytes": 20130}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-e-t070-p100 / seed 1201: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5", "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t070-p100/seed_1201/raw_samples.csv", "sha256": "ab42728b2bea5adc25950a81f569a3df5c94b8f6a1ebdc8f585edd4631df541d", "size_bytes": 19412}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-e-t085-p100 / seed 1200: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.85}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5", "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t085-p100/seed_1200/raw_samples.csv", "sha256": "06f0d2cfcf3215f70541d7c88c00cec89579b47d0975a8e99efdf35566240a8d", "size_bytes": 19118}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-e-t085-p100 / seed 1201: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.85}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3", "raw": {"relative_path":

"output/udlm/engineering_v5/v5-e-t085-p100/seed_1201/raw_samples.csv", "sha256": "327e7b74d31709973085a72743b62661abe2663f0969ea8ae44237ee6b1d88de", "size_bytes": 19054}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-e-t100-p100 / seed 1200: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t100-p100/seed_1200/raw_samples.csv", "sha256": "c5441a763ac9843d4a0865bbb0f33db9e47a9a1c55d2195de1bc29d1bbde9e51", "size_bytes": 18276}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-e-t100-p100 / seed 1201: {"checkpoint": "dce870e8d63453f73c428b9115f33556f67a21d45788d3112c2777ea82623005", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "f738b8b17de5c4704058018bbddacd7fed779c85248e33d199b68a648151e612", "prior_variant": "empirical_frequency", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-e-t100-p100/seed_1201/raw_samples.csv", "sha256": "dd15322742447031c8d04ef61c37f6e0a98f46f32c57ff4e50a258b32ef20693", "size_bytes": 18294}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-s-t050-p100 / seed 1200: {"checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-s-t050-p100/seed_1200/raw_samples.csv", "sha256": "67750d50a016eae3431aa95d3ec47a4976e1c6295323bc69777bcefa49d4caf9", "size_bytes": 21059}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-s-t050-p100 / seed 1201: {"checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-s-t050-p100/seed_1201/raw_samples.csv", "sha256": "dd15322742447031c8d04ef61c37f6e0a98f46f32c57ff4e50a258b32ef20693", "size_bytes": 18294}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

{ "active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": { "compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216, "physical_index": "3", "raw": { "relative_path": "output/udlm/engineering_v5/v5-s-t050-p100/seed_1201/raw_samples.csv", "sha256": "1de893cf2f402a2021f8eb08dd6e9eca5a55e3627ada5dc54b990c1e7fe04af2", "size_bytes": 21005, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05" }

v5-s-t070-p100 / seed 1200: { "checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": { "diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7, "gpu": { "cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": { "compute_mode": "Default", "compute_processes": { "pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880, "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "launch_policy": { "active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": { "compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216, "physical_index": "5", "raw": { "relative_path": "output/udlm/engineering_v5/v5-s-t070-p100/seed_1200/raw_samples.csv", "sha256": "ca74886bf6be3cee9a7a51f43724141fb63bd9ffd26b6ebc5a99c658d6cb8290", "size_bytes": 19956, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05" }

v5-s-t070-p100 / seed 1201: { "checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": { "diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7, "gpu": { "cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": { "compute_mode": "Default", "compute_processes": { "pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880, "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "launch_policy": { "active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": { "compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216, "physical_index": "3", "raw": { "relative_path": "output/udlm/engineering_v5/v5-s-t070-p100/seed_1201/raw_samples.csv", "sha256": "fc3769c7bcf5071f1b8984559bb86f1abac89b8c07b47828c0770b8d2081b298", "size_bytes": 19596, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05" }

v5-s-t085-p100 / seed 1200: { "checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": { "diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.85, "gpu": { "cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": { "compute_mode": "Default", "compute_processes": { "pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880, "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 5, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "launch_policy": { "active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than", "logical_device": { "compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216, "physical_index": "3", "raw": { "relative_path": "output/udlm/engineering_v5/v5-s-t085-p100/seed_1200/raw_samples.csv", "sha256": "db612b3231bc4f17e07b052863303215c6de0d5930e8a8114adced6f1ed89f7d", "size_bytes": 19991, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05" }

v5-s-t085-p100 / seed 1201: { "checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": { "diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0,

"softmax_temp": 0.85}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-s-t085-p100/seed_1201/raw_samples.csv", "sha256": "86c26ea024fb46143df0f5de9c67801411e5489c2a6db8dad7e37064c4e2b947", "size_bytes": 19489}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-s-t100-p100 / seed 1200: {"checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-s-t100-p100/seed_1200/raw_samples.csv", "sha256": "6356b28ba7eeea9c2128ac0add101b0705c8f408b2309963798fa71fb282be8e", "size_bytes": 18533}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-s-t100-p100 / seed 1201: {"checkpoint": "100f467b94766f2c87cc734398c8590bd14a1c8e0722ce31dbca7b446d986e72", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": "a5bbbe437cb732c516e3f2ad461b66de7f55e2e160e02c724f5e778cffb46fb7", "prior_variant": "schedule_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-s-t100-p100/seed_1201/raw_samples.csv", "sha256": "f56974a534d8a39435858a77c3c341c459780fb6f4dda83798bd4ab55430ae34", "size_bytes": 18519}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t050-p100 / seed 1200: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 9, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t050-p100/seed_1200/raw_samples.csv", "sha256": "f941d5b4cc21edeef012814945ec2ac534d0d5ab655336e93ef4ac1b821168e6", "size_bytes": 21718}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t050-p100 / seed 1201: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.5}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 9, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": 5}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t050-p100/seed_1201/raw_samples.csv", "sha256": "d27ee7308b3596eb46dae8c165aafa22ca96aeb21135a6411d7ab030c7ce2b28", "size_bytes": 20790}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t070-p100 / seed 1200: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": 5}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t070-p100/seed_1200/raw_samples.csv", "sha256": "bbbeb58d2bbb6faeece43f13fc7a21d2204137f61406eac9a48e7f352d42cd820", "size_bytes": 20237}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t070-p100 / seed 1201: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.7}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": 3}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t070-p100/seed_1201/raw_samples.csv", "sha256": "f9b92cd25a1cff261b49f516fc350bda523480b5320e59c8f2b259bcf3b0adc6", "size_bytes": 20340}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t085-p100 / seed 1200: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.85}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 8, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": 5}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t085-p100/seed_1200/raw_samples.csv", "sha256": "81d10fb2e12e415f8c54fb1471ddbd2a22242a90a9b18c2b8e672a0edc05ce72", "size_bytes": 19906}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t085-p100 / seed 1201: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 0.85}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 1284}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 1307, "name": "NVIDIA RTX A6000", "utilization_percent": 9, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t085-p100/seed_1201/raw_samples.csv", "sha256": "76cb87c29b6008868ad7a92ee99bc1d266030d4290f0b344185c8f99aa5f89b0", "size_bytes": 19645}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t100-p100 / seed 1200: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 275418, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 880}], "index": 5, "memory_total_mib": 49140, "memory_used_mib": 903, "name": "NVIDIA RTX A6000", "utilization_percent": 7, "uuid": "GPU-134e9a1b-cf80-5c0e-b0c4-048fff8e6d3c"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "5"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t100-p100/seed_1200/raw_samples.csv", "sha256": "a6a569384cc07e077edc60ce7e3883a0b4866a62d64024a6526dad5dea81dfa3", "size_bytes": 19414}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}

v5-r-t100-p100 / seed 1201: {"checkpoint": "d0310d2e2402043ab9ab6a99268263581a35d60fb2262ddd988ea38cd6f6395c", "config": {"diffusion_type": "udlm", "exclude_special_tokens": false, "inference_eps": 1e-05, "min_add_len": 40, "num_steps": 128, "prior_metadata_sha256": null, "prior_variant": "release_uniform", "randomness": 0.0, "raw_loo_top_p": 1.0, "softmax_temp": 1.0}, "gpu": {"cuda_visible_devices": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959", "final_probe": {"compute_mode": "Default", "compute_processes": [{"pid": 227850, "process_name": "/home/meir.roketlishvili/miniconda3/envs/SGPO_repro/bin/python", "used_memory_mib": 1284}], "index": 3, "memory_total_mib": 49140, "memory_used_mib": 1307, "name": "NVIDIA RTX A6000", "utilization_percent": 9, "uuid": "GPU-2cd1aa5b-616e-e6d3-b54d-49241cc8f959"}, "launch_policy": {"active_compute_processes_allowed": true, "inventory_scope": "all_nvidia_gpus", "max_utilization_percent": 10, "min_free_memory_mib": 30000, "requested_gpu_count": 2, "selection_method": "dynamic_idle_discovery", "utilization_comparison": "strictly_less_than"}, "logical_device": {"compute_capability": [8, 6], "logical_index": 0, "name": "NVIDIA RTX A6000", "total_memory_bytes": 50897289216}, "physical_index": "3"}, "raw": {"relative_path": "output/udlm/engineering_v5/v5-r-t100-p100/seed_1201/raw_samples.csv", "sha256": "a1bb88ced33cd8c2b6716b415fcfe2a83e54cd7d9eb2012d676be09ff54f3049", "size_bytes": 19966}, "source": "12e13ddf9b1323bbfb60fe289b5dd363ca0c3c05"}