UDLM resolution: paired diagnostic

Differences are 512 minus 128 predictor evaluations. Both objectives use temperature 1.0, identical checkpoint bytes, and 100 requests for each of seeds 17300 and 17301. 512 steps cost four times the model calls. Two seeds provide an engineering diagnostic, not a superiority claim.

Arm	Decode	Metric	Mean delta	Seed SD
CT	repaired	validity	-0.02000	+0.02828
CT	repaired	uniqueness	+0.00505	+0.00714
CT	repaired	quality	-0.01000	+0.05657
CT	repaired	diversity	-0.00073	+0.00302
CT	strict	validity	-0.01500	+0.04950
CT	strict	uniqueness	+0.00000	+0.00000
CT	strict	quality	-0.03000	+0.05657
CT	strict	diversity	-0.01071	+0.00112
CE	repaired	validity	-0.01000	+0.01414
CE	repaired	uniqueness	+0.00505	+0.00714
CE	repaired	quality	+0.03500	+0.02121
CE	repaired	diversity	+0.00485	+0.00419
CE	strict	validity	+0.04000	+0.05657
CE	strict	uniqueness	+0.00000	+0.00000
CE	strict	quality	+0.02000	+0.01414
CE	strict	diversity	+0.00624	+0.00202

The accompanying report.pdf contains configurations, checkpoint identities, per-run metrics, GPU mappings, timings, and the contextual MDLM comparator. report.json additionally records lexical error counts and both per-seed differences.